Giantbook

2015-02-03, rev. 8f1be59

Study the structure of social networks in the Erdős-Renyi model, such as the emergence of the giant component. This exercise serves as a very simple client for the Union–Find data structure.

Description

A social network consists of individuals and their connections; you are welcome to think of the structure of "friends" defined by Face-Book. In particular, the social networks studied in this exercise are symmetric in the sense that if p is connected to q then q is connected to p. (Other networks, such as the structure of "followers" on Twitter, are not symmetric.)

We are interested in the behaviour of these networks as they grow, *i.e.*, when connections are added. In particular, we want to know at which times the following events occur:

nonisolated The network is *nonisolated* when everybody has at least one connection.

giant component The network has a *giant* component when it contains a component with at least half the individuals.

connected The network is connected when everybody is connected to everybody else. In other words, when does the network consists of a single component that contains all individuals?

We study the dynamics of such a network in a very simple model called the *Erdős-Renyi model*, where in the beginning everybody has no connections, and individuals form links completely at random. Clearly, in the beginning, the network is neither nonisolated, connected, nor has a giant component. We want to know when the three events occur, and how the timing depends on N, the initial number of individuals.

Let me be very specific about the random model: In each round, create a connection between two individuals selected uniformly at random. It does not matter if these individuals have been selected in a previous round, of if they are already connected. (To make things easy, they can even be the same individual.)

Figure 1: A tiny evolving network. The *nonisolated* event happens at $T = 3$, the giant component emerges at $T = 4$, and the whole network becomes connected at $T = 6$.

Deliverables

1. Your implementation of `GiantBook.java`.

2. A modified Union–Find data structure in a separate file. Call this `MyUnionFind.java`.

3. A project report in PDF, you are strongly encouraged to use the skeleton on the last page.

It is important that the functionality of the union–find data structure and of the client (`GiantBook.java`) are clearly separated. In particular, don't put everything in one file. Your extensions of the functionality of the abstract data type for union–find should be in terms of meaningful, general methods that do not refer to the functionality of the client. [1]

Tips

The simplest event to simulate is *connected*, because the Union–Find structure in the book already tells you if there is just one component. Therefore, I suggest you completely solve the whole exercise for this event first, including filling out the relevant table column in the report. The book's standard library has useful methods on page 30. Make sure it works correctly for very small instances ($N = 10$) by following the computation step by step (print out what happens, check it by hand).

After that, study the emergence of the giant component. This involves modifying the Union–Find implementations in the book. I needed four lines of code.

Finally, study the nonisolated individuals. There is an obvious way of changing the Union–Find structure to detect the existence of isolated individuals (simply run through all elements and see if the size of their component is 1), but that will probably be too slow to be useful for large N. You need to do something that is both simpler and much faster. Again, I can do this in four lines of code.

Oh, and one more thing: it's probably a good idea to print the values using scientific notation, i.e., like `6.96e+02`. You can use the `StdOut.printf` method with format string "`%.2e`" for that.

Questions

My programme is slow. You should be able to detect the emergence of a giant component for $N = 10,000,000$ within seconds on a modern machine. If not, something is wrong. I am using weighted quick-union.

What about those standard deviations? This is standard scientific practice. When we perform computational experiments that involve randomness, the results depend on the random choices. Therefore we can't just run your experiment once and report the result. Instead, the experimented is repeated T times, and we report the

[1] For instance, `boolean hasGiantComponent()` is not a good method for your union–find data type, because it needs to know how a giant component is defined in this particular application. On the other hand, `int maxComponentSize()` would be a perfectly fine method for a union–find data type that could be useful to many other clients as well.

average (or "mean") of those experiments. So how large should T be? That's where the standard deviation comes in.[2] This is not a course in statistics, so we won't make a big deal out of this, but your standard deviation should be an order of magnitude smaller than the mean value you report. For example, if you report mean 536, your standard deviation should be at most something like 63, but not 251. Otherwise you need to increase T. For small experiments (say, $N = 10$) this is not important. The standard library has methods for both computing both the mean and standard deviation of a sequence of values `StdStats`, so you don't write any code for this.

[2] Read Wikipedia's entry on standard deviation for more.

Perspective

Want to learn more? You can read about the Erdős–Rényi model or more advanced models for social networks like the Watts–Strogatz model at Wikipedia. You can change your simulation to use that model instead, for example, and see what happens.

If you don't like the random graph models, check out the Stanford Network Analysis Project for real life data with large datasets. For example, running your code on the (famous) Enron email network should give the same value for the size of the largest component (called WCC in the dataset statistics table.)

Giantbook Report

by Alice Cooper and Bob Marley[3]

Results

The following table summarises our results. It shows the average number of random connections needed before the emergence of the giant component ("giant"), the disappearance of the last isolated individual ("no isolated"), and when the network becomes connected ("connect").

N	T	giant	(stddev)	no isolated	(stddev)	connected	(stddev)
100	100	7.15×10^1	5.77	2.68×10^2	6.32×10^1	2.69×10^2	6.3×10^1
1000							
10^4							
10^5							
10^6							
10^7							

Our main findings are the following: The first thing that happens is that [...] [4], which happens at a time [...] [5] in N. Perhaps surprisingly, two of the events seem to happen simultaneously, namely [...] and [...], which happen at time [...] in N.

Implementation details

We have based our union–find data type on $[\cdots]$.java from Sedgewick and Wayne: *Algorihthms, 4th ed.*, Addison–Wesley (2011). We added a method [...] by adding the following lines to [...]:[6]

Assuming we never run out of memory or heap space, if we would let our algorithm for detecting the emergence of a giant component run for 24 hours, it could compute the answer for $N = [\ldots]$.

We've run the code using a quick-find implementation as well. In 1 hour, we were able to handle and instance of size $N = [\ldots]$.

Discussion

We defined the giant component to have size at least αN for $\alpha = \frac{1}{2}$. [7] The choice of α is arbitrary, choosing other constants (such as $\alpha = \frac{1}{10}$ or $\alpha = \frac{9}{10}$) gave essentially the same results. The choice of constant is important; choosing [...] changes the experiment completely because [...].